


Today: Outline

- **ConvNets:** multiplication vs convolution; filters (or kernels); convolutional layers; 1D and 2D convolution; pooling layers; LeNet, CIFAR10Net
- **Reminders:** *Pre-lec Material 2,
due: Friday, Jun 4*


*Problem Set 1,
due: Friday, Jun 4*

Neural networks: recap



Learn parameters via gradient descent

$$\min_{\Theta} J(\Theta)$$




Backpropagation efficiently computes cost (forward pass) and gradient (backward pass)

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$$


Network architectures

Feed-forward


Fully connected




Recurrent



Convolutional



Face Recognition



[towardsdatascience.com]

Image Captioning



A young boy holding a baseball bat




A man riding a horse next to a building



Neural Networks

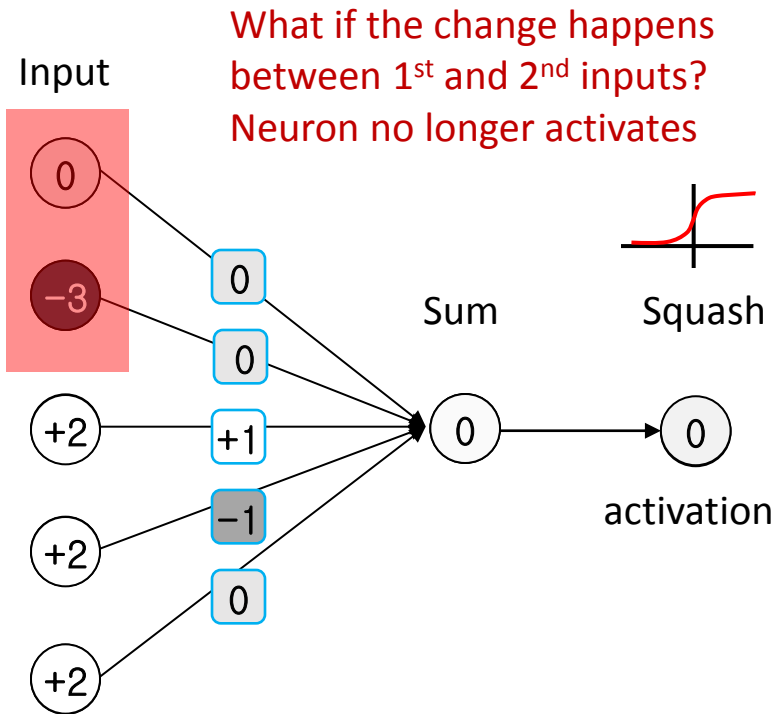
Convolutional Architectures

Multiplication vs convolution



- Recall, a neuron can be thought of as learning to spot certain features in the input
- E.g., this neuron detects change from high to low (light to dark) between 3rd and 4th inputs


Multiplication vs convolution



- Must have a new neuron for each new location of pattern???
- This is not efficient
- Solution: use convolution instead of multiplication

Multiplication vs convolution

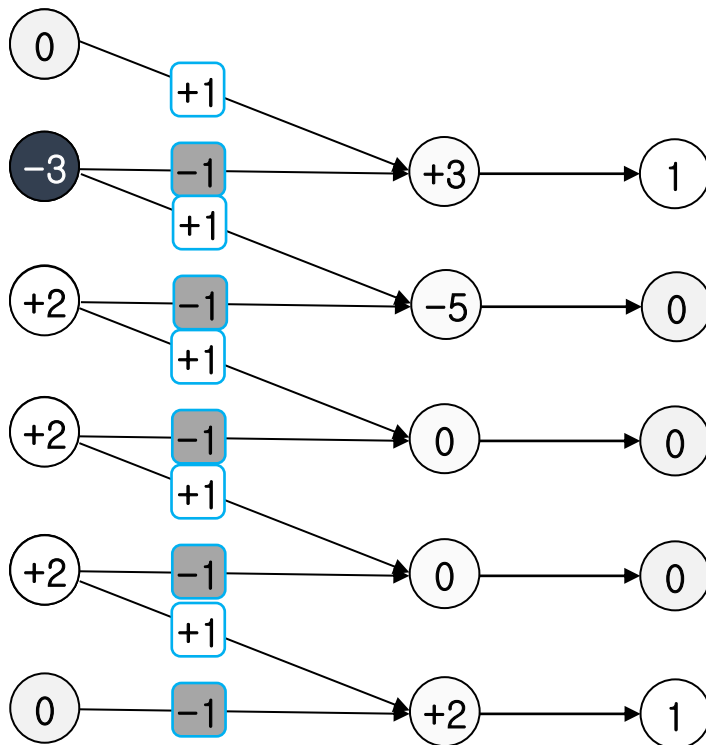
Input



- New weights are of size 2×1 ; called **filter**, or **kernel**
- New output is the size of input minus 1 because of **boundary**
- New convolutional neurons all **share** the same weights! This is much more efficient; we learn the weights once instead of many times for each position

Multiplication vs convolution


Padded
Input



- New output is the size of input minus 1 because of **boundary**
- We can fix the boundary effect by **padding** the input with 0 and adding one more neuron

Multiplication vs convolution


Padded
Input



- Note, we move the filter by 1 each time, this is called **stride**

Multiplication vs convolution


Padded
Input



- Note, we move the filter by 1 each time, this is called **stride**
- Stride can be larger, e.g. here is stride 2

Multiplication vs convolution

Padded
Input




To summarize, this layer has

- Input 5×1 , padded to 6×1
- Kernel 2×1 with weights $[+1, -1]$
- Stride 2
- Output 3×1

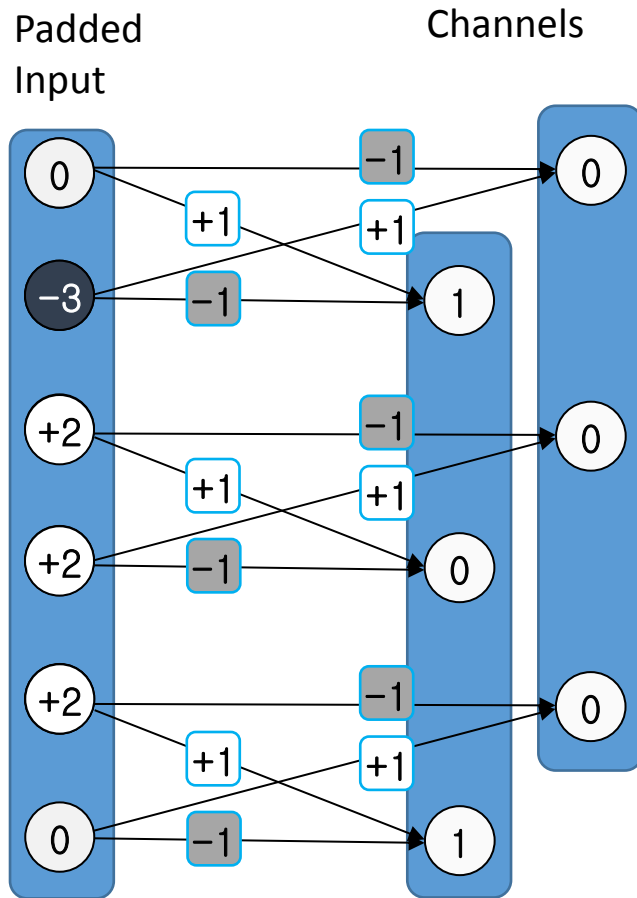
Multiplication vs convolution

Padded
Input




- We can add another filter, this time to detect opposite change with weights $[-1 \ 1]$
- Unique filters are called **channels**

Multiplication vs convolution



- We can add another filter, this time to detect opposite change with weights $[-1 \ +1]$
- Unique filters are called **channels**

Multiplication vs convolution



- We can add another filter, this time to detect opposite change with weights $[-1 \ 1]$
- Unique filters are called **channels**

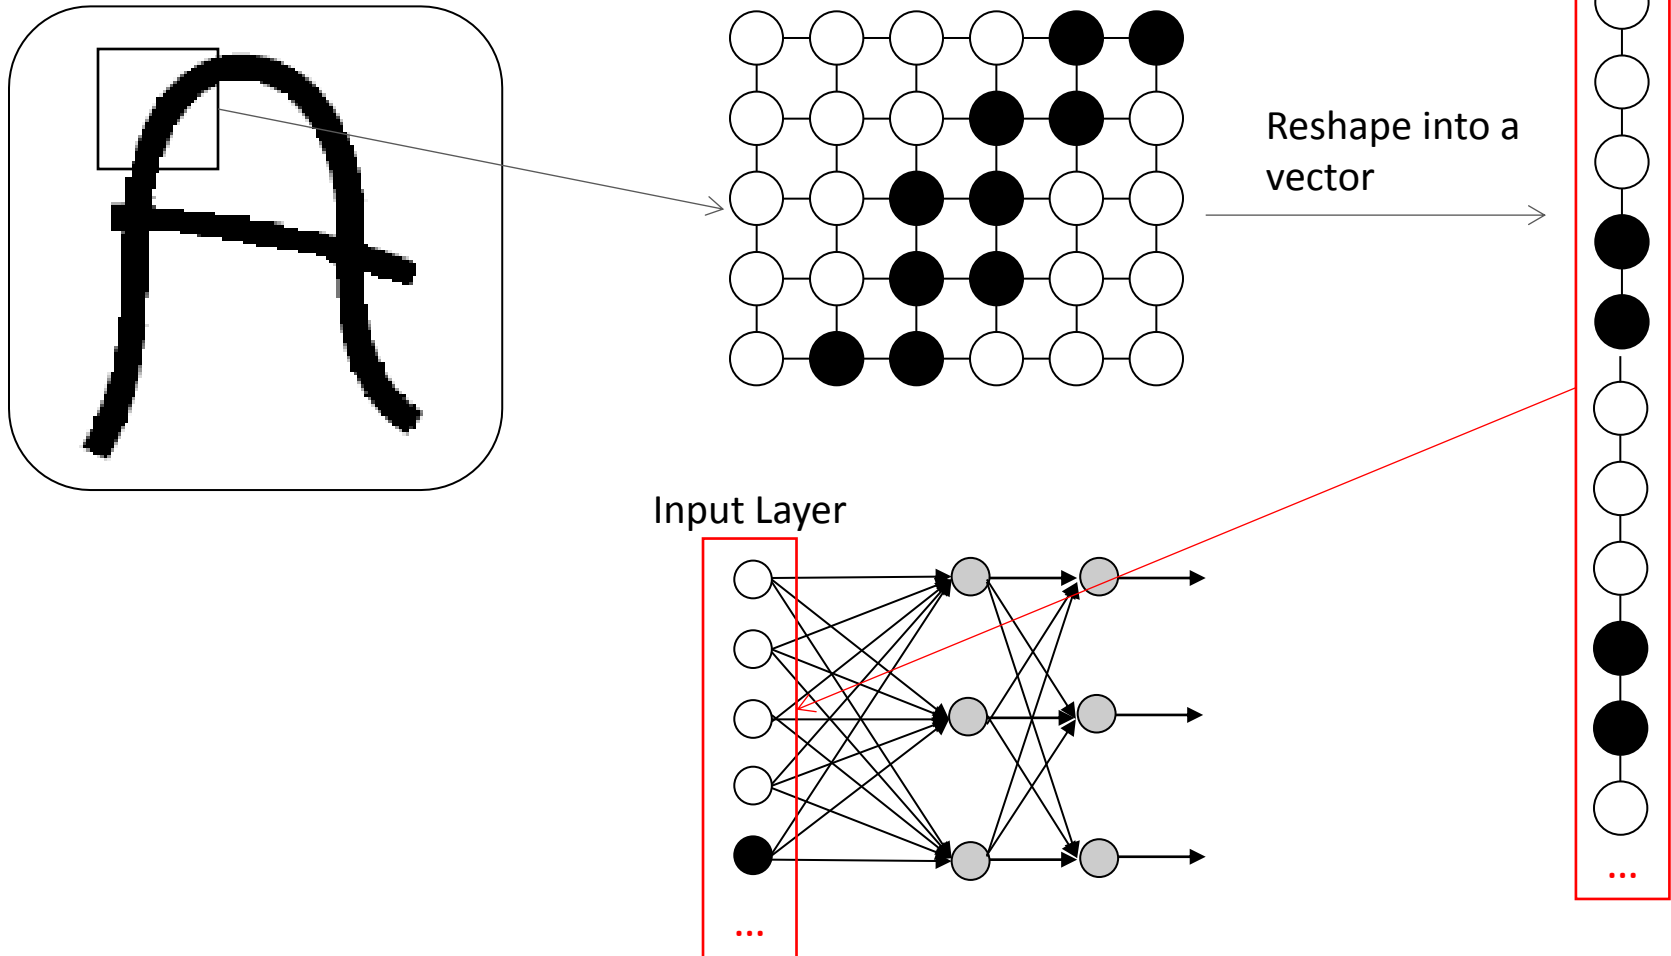


Convolutional Neural Networks

For images and other 2-D signals


Representing images

Fully connected




2D Input: fully connected network

Vectorize input by copying rows into a single column




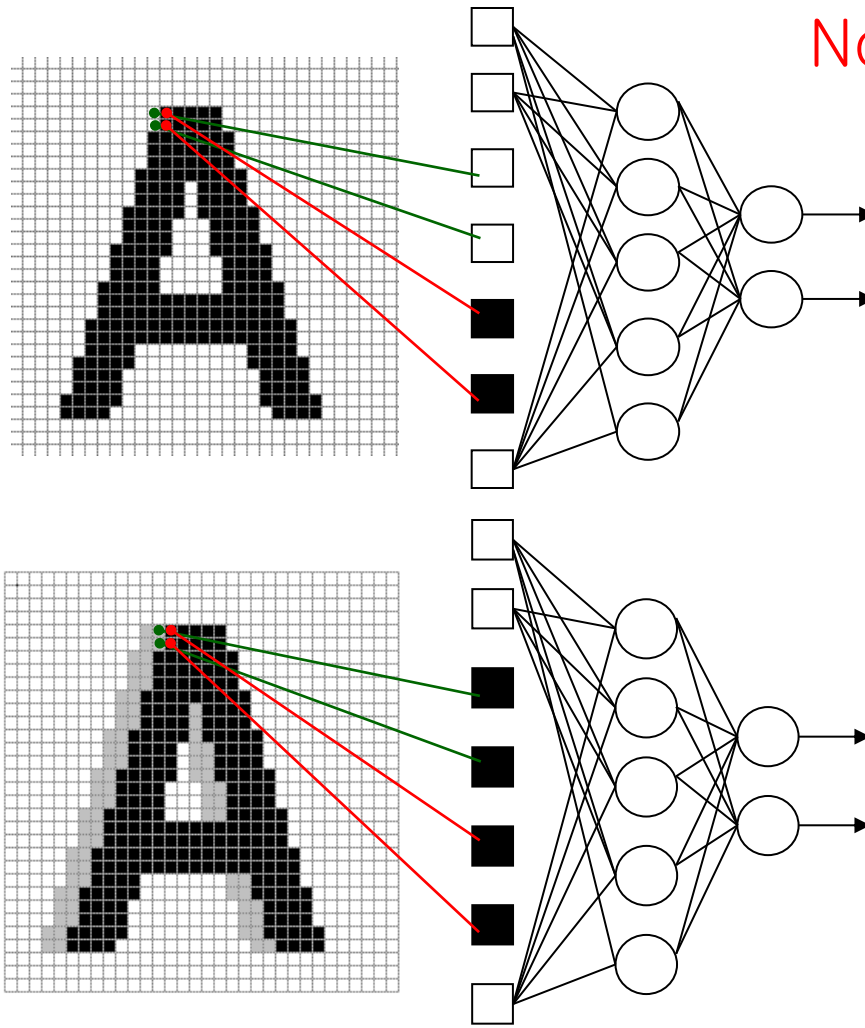
2D Input: fully connected network

Problem: shifting, scaling, and other distortion changes location of features



2D Input: fully connected network

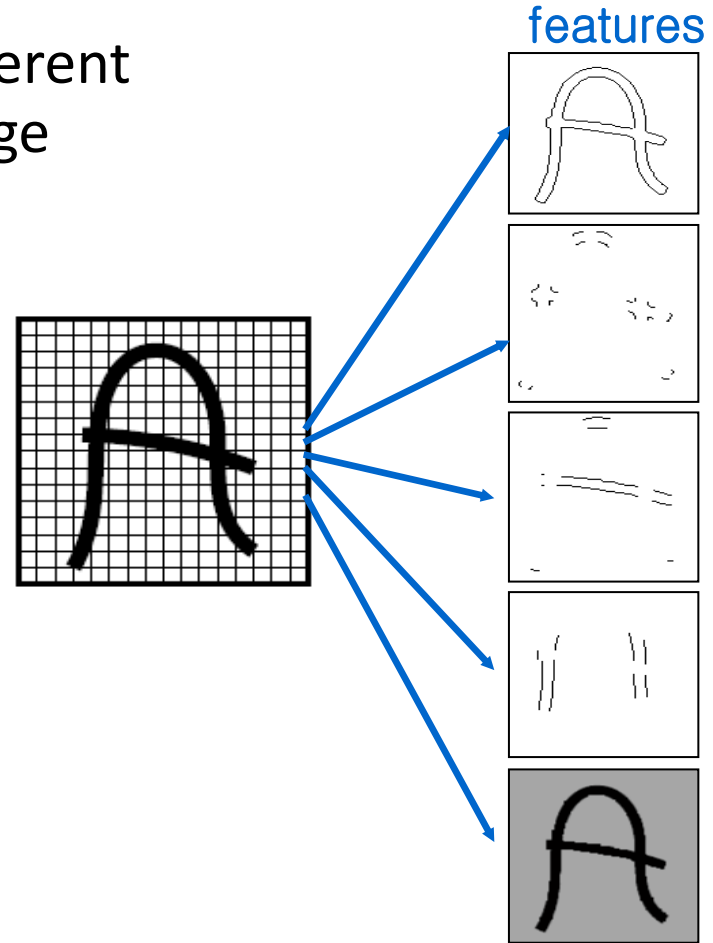
Not invariant to translation!





154 input change
from 2 shift left
77 : black to white
77 : white to black

Convolution layer in 2D



- detect the same feature at different positions in the input, e.g. image



Convolution layer in 2D



Convolution layer in 2D



$$a = f(w_{11}x_{11} + w_{12}x_{12} + w_{13}x_{13} + \cdots w_{33}x_{33})$$

What weights correspond to these output maps?



*These are output maps
before thresholding*

*Hint: filters look like the
input they fire on*





$$\frac{\partial f}{\partial x}(x, y)$$

$$\partial x$$



$$\frac{\partial f}{\partial y}(x, y)$$

$$\partial y$$



Where is Waldo?



Input



filter

What will the output map look like?



Input



filter

What will the output map look like?




filter

Output

Here is Waldo




Input



filter


Stacking convolutional layers

- Each layer outputs multi-channel **feature maps** (like images)
- Next layer learns filters on previous layer's feature maps




Pooling layers

- Convolution with stride > 1 reduces the size of the input
- Another way to downsize the feature map is with **pooling**
- A pooling layer subsamples the input in each sub-window
 - **max-pooling**: chose the max in a window
 - **mean-pooling**: take the average





Pooling layer

- the **pooling** layers reduce the spatial resolution of each feature map
- Goal is to get a certain degree of shift and distortion **invariance**




Pooling layer

- the **weight sharing** is also applied in pooling layers
- for mean/max pooling, no weights are needed

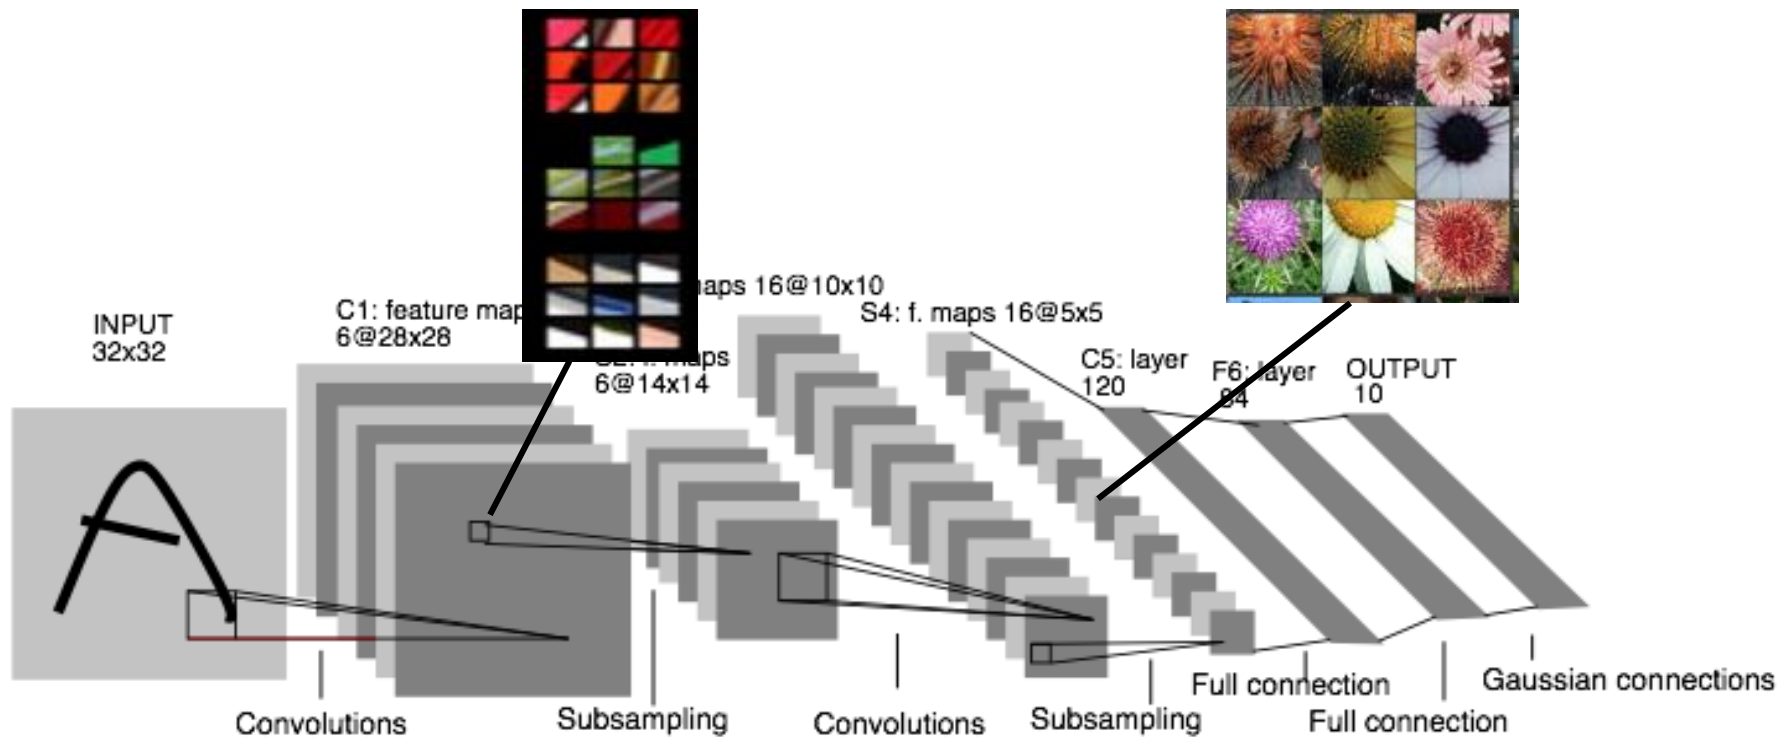


Putting it all together...



Convolutional Neural Network

A CNN is a better architecture for 2D signals



LeNet




Convolutional Neural Nets

Why they rule




Why CNNs rule: Translation invariance

- Output is invariant to translation of input
 - spatial translation for images
 - temporal translation for time sequences
- Note, not invariant to other transformations of input, such as large image rotation
- Pooling provides additional invariance to distortions



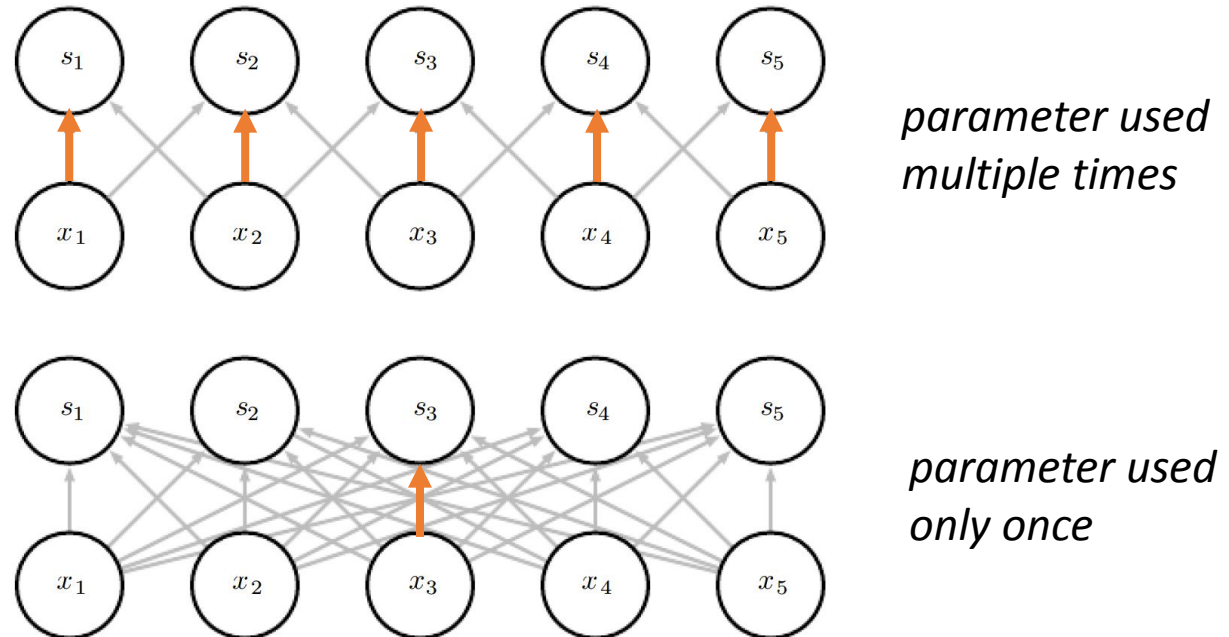
Why CNNs rule: Sparsity

- CNNs have sparse interactions, because the kernel is smaller than the input
- E.g. in thousands or millions pixel image, can detect small meaningful features such as edges
- Very efficient computation!
 - For m inputs and n outputs, matrix multiplication requires $O(m \times n)$ runtime (per example)
 - For k connections to each output, need only $O(k \times n)$ runtime
- Deep layers have larger effective inputs, or **receptive fields**



Why CNNs rule: Parameter sharing

- Kernel weights are shared across all locations
- Statistically efficient – learn from more data
- Memory efficient – store only k parameters, since $k \ll m$, this is much smaller than $m \times n$.



Alex Krizhevsky



Alex Krizhevsky

Dessa

Verified email at dessa.com

Machine Learning


 FOLLOW

TITLE	CITED BY	YEAR
Imagenet classification with deep convolutional neural networks A Krizhevsky, I Sutskever, GE Hinton Advances in neural information processing systems 25, 1097-1105	81373	2012
Dropout: a simple way to prevent neural networks from overfitting N Srivastava, G Hinton, A Krizhevsky, I Sutskever, R Salakhutdinov The journal of machine learning research 15 (1), 1929-1958	26336	2014
Learning multiple layers of features from tiny images A Krizhevsky, G Hinton		

Hence the name **AlexNet**



Why Now?



Neural Networks date back decades, so why the resurgence?

1. Big Data

- Larger Datasets
- Easier Collection & Storage

IMAGENET



2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable



3. Software

- Improved Techniques
- New Models
- Toolboxes








Convolutional Neural Nets

Example


CIFAR-10 Demo ConvJS Network



ReLU: rectified linear unit



ReLU function $g(x) = \max(0, x)$




input (32x32x3)



filter size 5x5x3, stride 1





filter size $5 \times 5 \times 3$, stride 1





input (32x32x3)



RELU

conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$






input (32x32x3)




filter size 5x5x3, stride 1



conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$





input (32x32x3)





pool (16x16x16)
pooling size 2x2, stride 2

filter size 5x5x3, stride 1



conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$





filter size $5 \times 5 \times 3$, stride 1




input ($32 \times 32 \times 3$)




pool ($16 \times 16 \times 16$)
pooling size 2×2 , stride 2


conv ($32 \times 32 \times 16$) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$




filter size $5 \times 5 \times 16$, stride 1




RELU



conv ($16 \times 16 \times 20$) params: $20 \times 5 \times 5 \times 16 + 20 = 8020$



pool ($8 \times 8 \times 20$)
pooling size 2×2 , stride 2



input (32x32x3)



One more conv+RELU+pool:

conv (8x8x20)

filter size 5x5x20, stride 1

relu (8x8x20)

pool (4x4x20)

pooling size 2x2, stride 2

fc (1x1x10); parameters: $10 \times 320 + 10 = 3210$



softmax (1x1x10)



Dog

car


Cat

⋮

Softmax

$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

\vec{z}	The input vector to the softmax function, made up of (z_0, \dots, z_K)
z_i	All the z_i values are the elements of the input vector to the softmax function, and they can take any real value, positive, zero or negative. For example a neural network could have output a vector such as $(-0.62, 8.12, 2.53)$, which is not a valid probability distribution, hence why the softmax would be necessary.
e^{z_i}	The standard exponential function is applied to each element of the input vector. This gives a positive value above 0, which will be very small if the input was negative, and very large if the input was large. However, it is still not fixed in the range $(0, 1)$ which is what is required of a probability.
$\sum_{j=1}^K e^{z_j}$	The term on the bottom of the formula is the normalization term. It ensures that all the output values of the function will sum to 1 and each be in the range $(0, 1)$, thus constituting a valid probability distribution.
K	The number of classes in the multi-class classifier.



input (32x32x3)



One more conv+RELU+pool:

conv (8x8x20)

filter size 5x5x20, stride 1

relu (8x8x20)

pool (4x4x20)

pooling size 2x2, stride 2

fc (1x1x10); parameters: $10 \times 320 + 10 = 3210$



softmax (1x1x10)



Dog

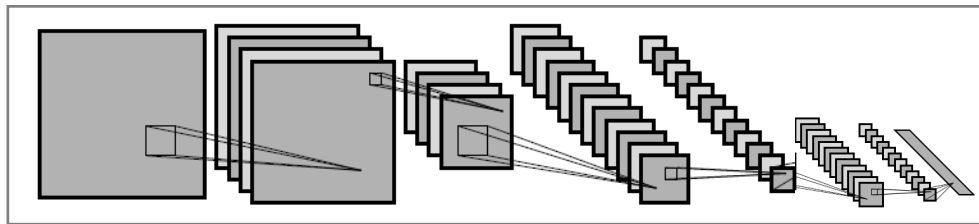
car

Cat

...


Testing the network

- Show top three most likely classes



<http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>

3D Convolutional Neural Networks



Application:

AI Generated Match Highlights

- IBM's produce the official match highlights of Wimbledon and US Open tennis tournaments.
- https://www.usopen.org/en_US/video/2017-08-31/1504233424.html

- Multi-modal System
- Bias Considerations

