

Today: Outline

- Homework Submission
- **ConvNets**: Example
- Training Strategies for Neural Networks
- Feature Extraction and Transfer Learning

- **Reminders**: *Pre-lec Material 2,*
due: Friday, Jun 4

Problem Set 1,
due: Friday, Jun 4

Anaconda Installation

- To run and solve assignments in this course, one must have a working IPython Notebook installation.
- The easiest way to set it up for both Windows and Linux is to:
 - install Anaconda: <https://www.anaconda.com/distribution/> (Python Version 3)
 - save and run this file to your computer
- If you are new to Python or its scientific library, Numpy, there are some nice Tutorials: <https://www.learnpython.org> and <http://scipy-lectures.org>
- In Windows after installation, search “Anaconda Navigator”. In the GUI menu, you can launch Jupyter Notebook or Jupyter Lab. These IDEs are based on a web-browser, you can enter localhost:8888 in a web browser and enter the work directory.

[Home](#)
[Environments](#)
[Learning](#)
[Community](#)

Applications on

base (root)

[Channels](#)
[Refresh](#)

JupyterLab
1.1.4

An extensible environment for interactive and reproducible computing, based on the Jupyter Notebook and Architecture.

[Launch](#)

Notebook
6.0.1

Web-based, interactive computing notebook environment. Edit and run human-readable docs while describing the data analysis.

[Launch](#)

Spyder
3.3.6

Scientific PYTHON Development Environment. Powerful Python IDE with advanced editing, interactive testing, debugging and introspection features

[Launch](#)

Glueviz
0.15.2

Multidimensional data visualization across files. Explore relationships within and among related datasets.

[Install](#)

Orange 3
3.23.0

Component based data mining framework. Data visualization and data analysis for novice and expert. Interactive workflows with a large toolbox.

[Install](#)

RStudio
1.1.456

A set of integrated tools designed to help you be more productive with R. Includes R essentials and notebooks.

[Install](#)

VS Code
1.42.0

Streamlined code editor with support for development operations like debugging, task running and version control.

[Install](#)
[Documentation](#)
[Developer Blog](#)


Jupyter Notebook

1. Numpy Tutorial

1.1 [5pt] Modify the cell below to return a 5x5 matrix of ones. Put some code there and press **Ctrl+Enter** to execute contents of the cell. You should see something like the output below. [1]

(<https://docs.scipy.org/doc/numpy-1.13.0/user/basics.creation.html#arrays-creation>) [2]

(<https://docs.scipy.org/doc/numpy-1.13.0/reference/routines.array-creation.html#routines-array-creation>)

```
In [3]: import numpy as np
import matplotlib.pyplot as plt

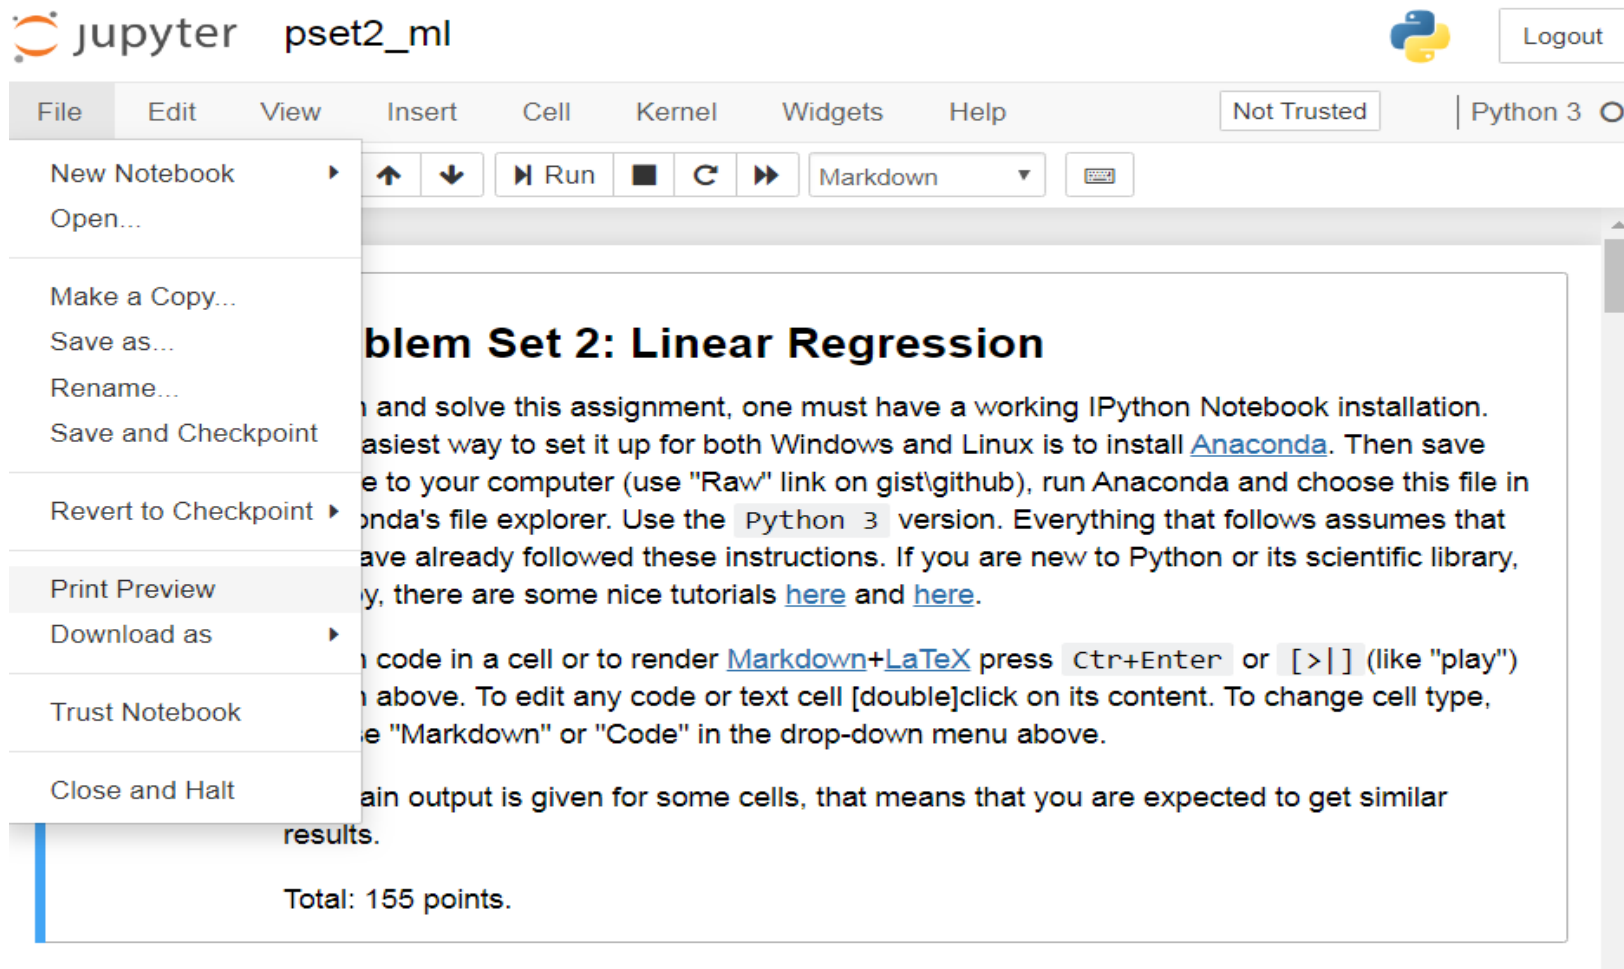
raise NotImplementedError("Replace this raise statement with the code "
                           "that prints 5x5 matrix of ones")
```

```
[[1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1.]]
```

1.2 [5pt] Vectorizing your code is very important to get results in a reasonable time. Let A be a 10x10 matrix and x be a 10-element column vector. Your friend writes the following code. How would you vectorize this code to run without any for loops? Compare execution speed for different values of n with `%timeit`

(<http://python.readthedocs.io/en/stable/interactive/magics.html#magic-timeit>).

Jupyter notebook - Save as pdf (i)



The screenshot shows a Jupyter Notebook interface. The top bar includes the Jupyter logo, the notebook name 'pset2_ml', a Python logo, and a 'Logout' button. Below this is a menu bar with 'File', 'Edit', 'View', 'Insert', 'Cell', 'Kernel', 'Widgets', and 'Help'. To the right of the menu bar are 'Not Trusted' and 'Python 3' buttons. The 'File' menu is open, displaying options: 'New Notebook', 'Open...', 'Make a Copy...', 'Save as...', 'Rename...', 'Save and Checkpoint', 'Revert to Checkpoint', 'Print Preview', 'Download as', 'Trust Notebook', and 'Close and Halt'. The 'Download as' option is highlighted. The notebook content area shows the title 'Problem Set 2: Linear Regression' and several paragraphs of text. The first paragraph discusses the requirements for solving the assignment, mentioning the need for a working IPython Notebook installation and the recommendation to use Anaconda. The second paragraph provides instructions on how to set up the environment, including using the 'Raw' link on GitHub and running Anaconda. The third paragraph explains how to execute code in a cell or render Markdown+LaTeX, mentioning the 'Ctrl+Enter' or '>' shortcuts. The fourth paragraph describes how to edit code or text cells and change their type. The fifth paragraph mentions that some cells have output, indicating that the user is expected to get similar results. The final line of the notebook content is 'Total: 155 points.'

jupyter pset2_ml

Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3

New Notebook
Open...

Make a Copy...
Save as...
Rename...
Save and Checkpoint

Revert to Checkpoint

Print Preview
Download as

Trust Notebook

Close and Halt

Problem Set 2: Linear Regression

and solve this assignment, one must have a working IPython Notebook installation. The easiest way to set it up for both Windows and Linux is to install [Anaconda](#). Then save it to your computer (use "Raw" link on [gist/github](#)), run Anaconda and choose this file in Anaconda's file explorer. Use the Python 3 version. Everything that follows assumes that you have already followed these instructions. If you are new to Python or its scientific library, there are some nice tutorials [here](#) and [here](#).

To execute code in a cell or to render [Markdown+LaTeX](#) press `Ctrl+Enter` or `[>]` (like "play") as shown above. To edit any code or text cell [double]click on its content. To change cell type, click on the "Markdown" or "Code" in the drop-down menu above.

When output is given for some cells, that means that you are expected to get similar results.

Total: 155 points.

Jupyter notebook - Save as pdf (ii)

2/18/2020pset2_mi

Problem Set 2: Linear Regression

To run and solve this assignment, one must have a working IPython Notebook installation. The easiest way to set it up for both Windows and Linux is to install [Anaconda](https://www.continuum.io/downloads) (<https://www.continuum.io/downloads>). Then save this file to your computer (use "Raw" link on gist.github), run Anaconda and choose this file in Anaconda's file explorer. Use the Python 3 version. Everything that follows assumes that you have already followed these instructions. If you are new to Python or its scientific library, Numpy, there are some nice tutorials [here](https://www.learnpython.org/) (<https://www.learnpython.org/>) and [here](http://www.scipy-lectures.org/) (<http://www.scipy-lectures.org/>).

To run code in a cell or to render [Markdown](https://en.wikipedia.org/wiki/Markdown) (<https://en.wikipedia.org/wiki/Markdown>)+[LaTeX](https://en.wikipedia.org/wiki/LaTeX) (<https://en.wikipedia.org/wiki/LaTeX>) press `Ctrl+Enter` or `[>]` (like "play") button above. To edit any code or text cell [double]click on its content. To change cell type, choose "Markdown" or "Code" in the drop-down menu above.

If certain output is given for some cells, that means that you are expected to get similar results.

Total: 155 points.

1. Numpy Tutorial

1.1 [5pt] Modify the cell below to return a 5x5 matrix of ones. Put some code there and press `Ctrl+Enter` to execute contents of the cell. You should see something like the output below. [1]
(<https://docs.scipy.org/doc/numpy-1.13.0/user/basics.creation.html#arrays-creation>) [2]
(<https://docs.scipy.org/doc/numpy-1.13.0/reference/routines.array-creation.html#routines-array-creation>)

```
In [3]: import numpy as np
import matplotlib.pyplot as plt

raise NotImplementedError("Replace this raise statement with the code "
                          "that prints 5x5 matrix of ones")
```

```
[[1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1.]
 [1. 1. 1. 1. 1.]]
```

1.2 [5pt] Vectorizing your code is very important to get results in a reasonable time. Let A be a 10x10 matrix and x be a 10-element column vector. Your friend writes the following code. How would you vectorize this code to run without any for loops? Compare execution speed for different values of n with `%timeit`
(<http://python.readthedocs.io/en/stable/interactive/magics.html#magic-timeit>).

Print14 pages

DestinationSave as PDF

PagesAll

LayoutPortrait

More settings

SaveCancel

Submission

CS 523

Summer 2021

Entry Code: **BP8RKK**

DESCRIPTION

Deep Learning

◆ ACTIVE ASSIGNMENTS

RELEASED

DUE (EDT) ▼

◆ SUBMISSIONS

Problem Set 1

MAY 28

JUN 04 AT 11:59PM


1




Convolutional Neural Nets

Example

CIFAR-10 Demo ConvJS Network






input (32x32x3)



filter size 5x5x3, stride 1





filter size $5 \times 5 \times 3$, stride 1





input ($32 \times 32 \times 3$)



RELU

conv ($32 \times 32 \times 16$) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$






input (32x32x3)




filter size 5x5x3, stride 1



conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$





input (32x32x3)





pool (16x16x16)
pooling size 2x2, stride 2

filter size 5x5x3, stride 1



conv (32x32x16) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$





filter size $5 \times 5 \times 3$, stride 1




input ($32 \times 32 \times 3$)




pool ($16 \times 16 \times 16$)
pooling size 2×2 , stride 2


conv ($32 \times 32 \times 16$) params: $16 \times 5 \times 5 \times 3 + 16 = 1216$




filter size $5 \times 5 \times 16$, stride 1




RELU



conv ($16 \times 16 \times 20$) params: $20 \times 5 \times 5 \times 16 + 20 = 8020$



pool ($8 \times 8 \times 20$)
pooling size 2×2 , stride 2



input (32x32x3)



One more conv+RELU+pool:

conv (8x8x20)

filter size 5x5x20, stride 1

relu (8x8x20)

pool (4x4x20)

pooling size 2x2, stride 2

fc (1x1x10); parameters: $10 \times 320 + 10 = 3210$



softmax (1x1x10)



Dog

car


Cat

⋮

Softmax

$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

\vec{z}	The input vector to the softmax function, made up of (z_0, \dots, z_K)
z_i	All the z_i values are the elements of the input vector to the softmax function, and they can take any real value, positive, zero or negative. For example a neural network could have output a vector such as $(-0.62, 8.12, 2.53)$, which is not a valid probability distribution, hence why the softmax would be necessary.
e^{z_i}	The standard exponential function is applied to each element of the input vector. This gives a positive value above 0, which will be very small if the input was negative, and very large if the input was large. However, it is still not fixed in the range $(0, 1)$ which is what is required of a probability.
$\sum_{j=1}^K e^{z_j}$	The term on the bottom of the formula is the normalization term. It ensures that all the output values of the function will sum to 1 and each be in the range $(0, 1)$, thus constituting a valid probability distribution.
K	The number of classes in the multi-class classifier.



input (32x32x3)



One more conv+RELU+pool:

conv (8x8x20)

filter size 5x5x20, stride 1

relu (8x8x20)

pool (4x4x20)

pooling size 2x2, stride 2

fc (1x1x10); parameters: $10 \times 320 + 10 = 3210$



softmax (1x1x10)



Dog


car

Cat

⋮


Testing the network

- Show top three most likely classes



<http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>

3D Convolutional Neural Networks

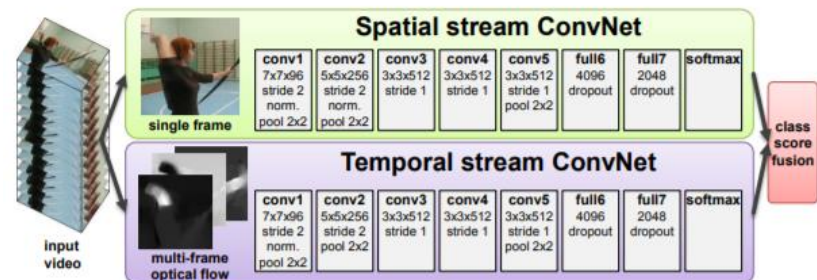


Application:

AI Generated Match Highlights

- IBM's produce the official match highlights of Wimbledon and US Open tennis tournaments.
- https://www.usopen.org/en_US/video/2017-08-31/1504233424.html

- Multi-modal System
- Bias Considerations





Neural Networks

Training Strategies

Universality

- Why study neural networks in general?
 - Neural network can approximate any continuous function!
 - <http://neuralnetworksanddeeplearning.com/chap4.html>

Why Study Deep Networks?





Figure 6.6: Empirical results showing that deeper networks generalize better when used to transcribe multi-digit numbers from photographs of addresses. Data from [Goodfellow et al. \(2014d\)](#). The test set accuracy consistently increases with increasing depth. See figure 6.7 for a control experiment demonstrating that other increases to the model size do not yield the same effect.

Efficiency of convnets



Activation Functions

- ReLU: $g(x) = \max(0, x)$
- Leaky ReLU: $g(x) = \max(0, x) + \alpha \min(0, x)$ ($\alpha \approx .01$)
- Tanh: $g(x) = 2\sigma(2x) - 1$
- Radial Basis Functions: $g(x) = \exp(-(w - x)^2 / \sigma^2)$
- Softplus: $g(x) = \log(1 + e^x)$
- Hard Tanh: $g(x) = \max(-1, \min(1, x))$
- Maxout: $g(x) = \max_{j \in \mathbb{G}} x_j$
-

Architectures

- Some commonly referred to architectures:
 - AlexNet
 - VGG16/19
 - GoogleNet
 - ResNet
 - WideResNet
 - Inception
 - ...

Architecture Design and Training Issues

- How many layers? How many hidden units per layer? How to connect layers together? How to optimize?
 - Cost functions
 - L2/L1 regularization
 - Data Set Augmentation
 - Early Stopping
 - Dropout
 - Minibatch Training
 - Momentum
 - Initialization
 - Batch Normalization

Cost Functions

- For regression problems, quadratic error is typical
- For classification, one typically uses softmax outputs with cross-entropy error function

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$


Regularization

- In machine learning, we care about *generalization performance*, not just training error
- With many parameters, models are prone to *overfitting*
- How to regularize?
 - Restrictions on parameter values
 - Adding terms to the objective function
 - Examples: L2 or L1 regularization
- Other ways to improve generalization?

Data Set Augmentation

- The more data, the better for generalization (usually)
- Sometimes we can augment our existing data set
 - Example: for image classification, mirror-image all images to double the size of the training set
 - Injecting noise to training data is also a form of data augmentation

Early Stopping



Dropout




Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

Dropout





Figure 4: Test error for different architectures with and without dropout. The networks have 2 to 4 hidden layers each with 1024 to 2048 units.

Architecture Design and Training Issues

- How many layers? How many hidden units per layer? How to connect layers together? How to optimize?
 - Cost functions
 - *L2/L1 regularization*
 - *Data Set Augmentation*
 - *Early Stopping*
 - *Dropout*
 - Minibatch Training
 - Momentum
 - Initialization
 - Batch Normalization
- 
- Avoid Overfitting*

Minibatch Training

- Gradient descent uses all training points, fully online (stochastic) methods update using a single point
- Many deep learning methods fall in between
- Example: computing the mean of a set of samples
 - Standard error based on a sample of n points is σ / \sqrt{n}
 - Consider using 100 versus 10,000 samples
 - Latter requires 100x more computation but reduces error by factor of 10

Minibatch Training

- Larger batches provide a more accurate estimate of the gradient. If all examples in the batch are processed in parallel, amount of memory scales with batch size.
- Small batches can offer a regularizing effect due to the noise added during the learning process.
- When using GPUs it is common for power of 2 batch sizes to offer better runtime; typical sizes range from 32 to 256, with 16 being common for large models.
- Minibatches should be selected randomly!

GPUs

- NVIDIA TITAN V GPU




Mini-batches

- Gradients could be updated using:
 - One data point (too inaccurate)
 - All data points (too expensive)
 - Mini-batch (a good trade-off)
- The size of the mini-batch depends on:
 - How good of an approximation you need
 - How much GPU memory you have per GPU
 - How many GPUs you have
- GPUs can compute gradients of mini-batches in parallel, *i.e. Training on multiple GPUs: Divide and Conquer.*

Momentum

- Accumulates an exponentially decaying moving average of past gradients and continues to move in their direction
- exponentially weighed averages can provide us a better estimate which is closer to the actual derivate than our noisy batch calculations





Initialization

- Important: need to “break symmetry”
 - E.g. Choose weights randomly
- Combined with early stopping, can think of initialization as a prior on the weights
- Usually use uniform or Gaussian weights with a zero-mean
- Examples with m inputs and n outputs:

$$W_{ij} \sim U\left(-\frac{1}{\sqrt{m}}, \frac{1}{\sqrt{m}}\right) \quad W_{ij} \sim U\left(-\frac{6}{\sqrt{m+n}}, \frac{6}{\sqrt{m+n}}\right)$$

Batch Normalization

- High-level idea: whitening the data at each layer makes training faster
- Left: MNIST, Right: ImageNet



Data Independence

- NN models converging to the correct solution depends on the iid assumption
- i.e. that our data are independent and identically distributed
- Important to randomly shuffle examples!
Otherwise net can fail to converge